

Postgres Advanced Data Types

Fanavaran Anisa
Iran Linux House

Linux & Open Source Training Center

www.anisa.co.ir



Table of Contents

- JSON/JSONB
- Arrays
- HSTORE

JSON Data Type

- JSON data **types are for storing JSON** (JavaScript Object Notation) data, as specified in RFC 7159.
- Such data can also be stored as text, but the JSON data types have the advantage of **enforcing that each stored value is valid** according to the JSON rules. There are also assorted **JSON-specific functions and operators** available for data stored in these data types

PostgreSQL offers two types for storing JSON data:

- **JSON**
- **JSONB**

PostgreSQL also provides the **jsonpath** data type

JSON vs JSONB

- The json and jsonb data types accept almost identical sets of values as input.
- The major practical difference is **efficiency**.
 - The **json** data type stores an exact copy of the input text, which processing functions must **reparse** on each execution
 - **jsonb** data is stored in a decomposed binary format
 - slower input
 - significantly faster to process

Jsonb also **supports indexing**, which can be a significant advantage

Which One to Use?

Why JSONB is the Always Better ?

- jsonb does not preserve white space
- does not preserve the order of object keys
- does not keep duplicate object keys. If duplicate keys are specified in the input, only the last value is kept.

In general, **most applications should prefer to store JSON data** as jsonb, unless there are quite specialized needs, such as legacy assumptions about ordering of object keys.

Which One to Use?

Why JSONB is the Always Better ?

- jsonb does not preserve white space
- does not preserve the order of object keys
- does not keep duplicate object keys. If duplicate keys are specified in the input, only the last value is kept.

In general, **most applications should prefer to store JSON data** as jsonb, unless there are quite specialized needs, such as legacy assumptions about ordering of object keys.

Which One to Use?

Why JSONB is the Always Better ?

- jsonb does not preserve white space
- does not preserve the order of object keys
- does not keep duplicate object keys. If duplicate keys are specified in the input, only the last value is kept.

In general, **most applications should prefer to store JSON data** as jsonb, unless there are quite specialized needs, such as legacy assumptions about ordering of object keys.

```
CREATE TABLE journal (  
  id Int NOT NULL PRIMARY KEY,  
  day VARCHAR(10),  
  diary_information JSONB  
);
```

First Steps

```
INSERT INTO journal (id, day, diary_information)
VALUES ( 1, "Tuesday",
'{"title": "My first day at work", "Feeling": "Mixed feeling"}' );
```

```
SELECT * FROM journal
```

	id [PK] character varying	age integer	diary_information jsonb
1	2	27	{"title": "My first day at work", "Feeling": "Mixed feelin..."}

First Steps – Functions & Operators

->: This operator allows you to extract a specific value from a JSON object, you specify the key as a “child” to the “parent”.

```
SELECT Id, day,  
diary_information -> 'Feeling' AS Feeling  
FROM journal;
```

First Steps – Functions & Operators

->>: This operator allows you to extract a JSON object field as text without the quotes around it from a JSON object.

```
SELECT id, day, dairy_information ->> 'Feeling' as Feeling  
FROM products;
```

First Steps – Functions & Operators

json_agg: This function aggregates JSON values into a JSON array.

For example, `SELECT json_agg(my_column) FROM my_table;` will return a JSON array containing the values in the "my_column" column of the "my_table" table.

jsonb_set: This function updates a JSON object field with a new value

```
UPDATE journal SET diary_information =  
  jsonb_set( diary_information, '{Feeling}', '"Excited"' )  
WHERE id = 1;
```

First Steps – Functions & Operators

You can use a JSONB_BUILD_OBJECT function to insert a plain text record and this will convert it to JSON data

```
INSERT INTO journal (id, day, feeling)
VALUES ( 2, 'Wednesday',
JSONB_BUILD_OBJECT(
'Morning',
'Everybody is annoying today',
'Evening',
'Cannot wait to go home' )
);
```

@> Containment Operator

the @> operator tests if the left-hand JSON value contains the right-hand JSON value.

The <@ operator tests if the right-hand JSON value contains the left-hand JSON value

```
INSERT INTO employees (id, name, skills) VALUES  
( 1, 'John', '[{"name": "Python", "level": "Intermediate"},  
{"name": "JavaScript", "level": "Expert"}]' );
```

```
SELECT name FROM employees WHERE skills @ > '[{"name": "Python"}]' :: jsonb
```

@> Containment Operator

the @> operator tests if the left-hand JSON value contains the right-hand JSON value.

The <@ operator tests if the right-hand JSON value contains the left-hand JSON value

```
INSERT INTO employees (id, name, skills) VALUES
( 1, 'John', '[{"name": "Python", "level": "Intermediate"},
{"name": "JavaScript", "level": "Expert"}]' );
```

```
SELECT name FROM employees WHERE skills @ > '[{"name": "Python"}]' :: jsonb
```

JSON Path

```
SELECT jsonb_extract_path('{"brand": "Honda", "sold":  
false} '::jsonb, 'brand');
```

```
SELECT '{"brand": "Honda", "sold": false}'::jsonb @> '{"brand":  
"Honda"}'::jsonb AS result
```

```
SELECT  
'{"brand": "Honda", "sold": false}' :jsonb @@ '$.brand ==  
"Honda"' AS result_with_@@,  
'{"brand": "Honda", "sold": false}'::jsonb @? '$.brand ? (@ ==  
"Honda")' AS result_with_@?;
```

HSTORE

The hstore module implements the hstore data type for storing key-value pairs in a single value. The keys and values are text strings only.

```
CREATE EXTENSION hstore;
```

```
CREATE TABLE books ( id serial primary key, title  
VARCHAR (255), attr hstore );
```

```
INSERT INTO books (title, attr) VALUES (  
'PostgreSQL Tutorial',  
"paperback" => "243",  
"publisher" => "postgresqltutorial.com",  
"language" => "English",  
"ISBN-13" => "978-1449370000",  
"weight" => "11.2 ounces" );
```


HSTORE: Query/Update/Delete

```
SELECT
    title, attr -> 'weight' AS weight
FROM
    books
WHERE
    attr -> 'ISBN-13' = '978-1449370000';
```

```
UPDATE books
SET attr = attr || '"freeshipping"=>"yes"' :: hstore;
```

```
UPDATE books
SET attr = attr || '"freeshipping"=>"no"' :: hstore;
```

```
UPDATE books
SET attr = attr - 'weight'
WHERE title = 'PostgreSQL Cheat Sheet';
```

HSTORE – Advanced Queries

```
SELECT
  title,
  attr->'publisher' as publisher,
  attr
FROM
  books
WHERE
  attr ? 'publisher';
```

returns all rows with attr contains key publisher.

HSTORE - HSTORE – Advanced Queries

```
SELECT title
FROM books WHERE
attr @> '"weight"=>"11.2 ounces"' :: hstore;
```

```
SELECT
    title
FROM
    books
WHERE
    attr ?& ARRAY [ 'language', 'weight' ];
```

ARRAY

In PostgreSQL, an array of a collection of elements that have the same data type.

Every data type has its companion array type e.g., integer has an integer[] array type, character has character[] array type.

```
CREATE TABLE contacts (  
  id SERIAL PRIMARY KEY,  
  name VARCHAR (100),  
  phones TEXT []  
);
```

column_name data_type [][]

ARRAY – Insert Data

```
INSERT INTO contacts (name, phones)
VALUES('John Doe', ARRAY [ '(408)-589-5846', '(408)-589-5555' ] );
```

```
INSERT INTO contacts (name, phones)
VALUES('Lily Bush', '{ "(408)-589-5841" }'),
      ('William Gate', '{ "(408)-589-5842", "(408)-589-58423" }');
```

ARRAY – Query

```
SELECT
  name,
  phones [ 1 ]
FROM
  contacts;
```

```
SELECT
  name,
  phones
FROM
  contacts
WHERE
  '(408)-589-5555' = ANY (phones);
```

Output:

name	phones
John Doe	{(408)-589-5846, (408)-589-5555}

(1 row)

ARRAY

```
UPDATE contacts
SET phones [2] = '(408)-589-5843'
WHERE ID = 3
RETURNING *;
```

Output:

id	name	phones
3	William Gate	{(408)-589-5842, (408)-589-5843}

(1 row)

phones = '{"(408)-589-5843"}'

Expanding ARRAY

```
SELECT
  name,
  unnest(phones)
FROM
  contacts;
```

Output:

name		unnest
John Doe		(408)-589-5846
John Doe		(408)-589-5555
Lily Bush		(408)-589-5841
William Gate		(408)-589-5843

(4 rows)